

Below is a step-by-step **Node.js embedding guide for ChatKit**, written for people who **already built an agent workflow in OpenAI Agent Builder** and want to **embed that published workflow into a Node/Express website** (the “Nikon demo” pattern).

This guide assumes you downloaded the demo from GitHub: https://github.com/BitBloom-HQ/openai_agent_builder_nikondemo (the repo structure matches your ZIP), and that you already have a **published Workflow ID** from Agent Builder.

0) Who this guide is for

This guide is **not** about building the agent workflow logic in Agent Builder. It’s for the next step: **embedding your existing agent workflow into a Node.js site using ChatKit**, with a secure server-side session endpoint and a minimal browser-side script.

1) What you get in this demo project (folder map)

From the ZIP (and the GitHub repo), the important files are:

- `server.js`
Express server that exposes **POST /api/chatkit/session** and creates a ChatKit session **server-side** (so your API key never goes to the browser).
- `package.json`
Contains the libraries you install (Express, EJS, dotenv, axios).
- `views/index.ejs`
Main page template. Loads ChatKit from the CDN and includes the chat container partial.
- `views/partials/chat-container.ejs`
The floating chat panel markup (container + “Connecting...” status).
- `public/app-chatkit.js`
Browser logic: fetches a short-lived client secret from your server and mounts `<openai-chatkit>`.
- `public/chat.css`
Styling for the floating chat window.
- `public/styles.css (+ optional page CSS)`
Site styling (not required for ChatKit to work).

2) Prerequisites from Agent Builder (you already did this)

A) You must publish your workflow

Publishing gives you the **Workflow ID** you'll use in the Node server.

B) You must have credits / billing set up

If your account has no funds, the UI may look fine but responses will fail; the real clue is in Logs ("quota / billing").

C) Domain allowlist + public key (important, even for localhost)

In your OpenAI dashboard, whitelist the domain under **Security → Domain allowlist**, then use the **Public Key** for that domain. For local dev you can use a temporary domain (example mentioned: ngrok).

3) Install dependencies (from package.json)

From the extracted folder:

```
npm install
```

This project installs (from package.json in your ZIP):

- express
 - ejs
 - dotenv
 - axios
-

4) Create your .env file (keys live on the server)

Create a .env file in the project root (next to server.js).

You need **three** variables:

1. OPENAI_API_KEY

Your secret API key. **Server only, never in the browser.**

2. WORKFLOW_ID

Your published workflow id from Agent Builder.

3. CHATKIT_DOMAIN_PK

The **public key** you got from **Security → Domain allowlist** (used by the ChatKit CDN script in views/index.ejs).

Example:

OPENAI_API_KEY=sk-...

WORKFLOW_ID=wf_...

CHATKIT_DOMAIN_PK=pk_...

PORT=3000

Why this matters: the demo's "golden rule" is exactly this: don't ship your OpenAI key to the browser.

5) Run the demo site

Start it:

npm start

(or node server.js, same thing)

Open:

- http://localhost:3000

Then click "**Chat with the Agent**" and you should see the floating chat window and a "**Connecting to agent...**" status line while it fetches a client secret.

6) How the integration works (the real wiring)

A) The server endpoint: POST /api/chatkit/session

Your browser never talks to OpenAI directly for sessions. Instead:

browser → your server → OpenAI → client_secret → browser

That's the handshake pattern.

In this codebase, server.js uses **axios** (not the OpenAI SDK) to call:

- POST `https://api.openai.com/v1/chatkit/sessions`
- With headers:
 - Authorization: Bearer \${OPENAI_API_KEY}
 - OpenAI-Beta: chatkit_beta=v1
- And body:
 - workflow: { id: WORKFLOW_ID }
 - user: <some identifier>

The endpoint responds to the browser with:

- `client_secret`
- `expires_at`

It also normalizes common error types (quota, rate limit, unauthorized) so the frontend can show a friendly message.

B) The client script: `public/app-chatkit.js`

This file does 3 core jobs:

1. Calls your server:
 - `fetch('/api/chatkit/session', { method: 'POST' })`
2. Creates and mounts the web component:
 - `document.createElement('openai-chatkit')`
3. Injects configuration via:
 - `el.setOptions({ api: { getClientSecret() { ... } }, ...buildOptions() })`

So the ChatKit widget gets its session secret **on demand** from your backend.

C) The HTML: ChatKit CDN + mount point

In `views/index.ejs`, ChatKit is loaded with the public key:

- https://cdn.platform.openai.com/deployments/chatkit/chatkit.js?pk=<%= process.env.CHATKIT_DOMAIN_PK %>

And the chat UI mounts into the container from:

- views/partials/chat-container.ejs (the <div id="my-chat"> target + status line)
-

7) Customizing the experience (colors + greeting + placeholder)

You have two layers of customization:

A) ChatKit theme + text (in public/app-chatkit.js)

The demo's buildOptions() returns:

- theme.colorScheme: 'dark'
- theme.color.accent.primary: '#FFD400' (Nikon yellow)
- composer.placeholder: 'Ask about the Nikon D700...'
- startScreen.greeting: 'What can I help you with?'
- history.enabled: true
- locale: 'en-US'

To customize:

- Change accent color: edit theme.color.accent.primary
- Change the input placeholder: edit composer.placeholder
- Change the opening greeting: edit startScreen.greeting

B) Widget container styling (in public/chat.css)

This controls the floating panel itself:

- size (width/height)
- position (bottom/right)
- border, shadow, radius
- responsive behavior for mobile

If you want a bigger widget, update:

- `#chat-container { width: 360px; height: 640px; }`
-

8) Common issues (quick fixes)

“It shows the chat but doesn’t answer”

Most common causes:

- Wrong WORKFLOW_ID
- Wrong OPENAI_API_KEY
- Your endpoint path isn’t exactly /api/chatkit/session
- No billing/credits (check logs: quota/billing issue).

“Forbidden / domain issues”

Make sure you added your domain to **Security → Domain allowlist** and used the matching public key, even for local testing (use a temporary domain if needed).

9) Production safety rules (do not skip)

- Keep OPENAI_API_KEY on the server only.
- Treat the ChatKit client_secret as short-lived and fetch fresh when needed.
- Domain allowlist your production domain and use the correct public key.